

Operating Systems

CSCI 5806

Spring Semester 2021 — CRN 21176

Term Project — Step 0 — Preparation

Target completion date: Monday, January 18, 2021

Goals

- Learn about the five basic low-level UNIX file access functions
- Create two functions to display memory regions

Details

UNIX has five basic file access system calls:

- **int open(char *fn,int mode)**
Open the file whose name is given. The filename can be either a relative or absolute path. Mode indicates whether the file is opened read-only (**O_RDONLY**), write-only (**O_WRONLY**) or read-write (**O_RDWR**). The function returns a nonnegative integer — the *file handle* — if the function succeeds, or **-1** if the function fails.
- **void close(int fd)**
Close the file whose handle is given.
- **ssize_t read(int fd,void *buf,size_t count)**
Reads the given number of bytes from the given open file, placing the bytes in the given buffer. The location of the first byte read is given by the file's *cursor*; subsequent bytes are read sequentially and the cursor is advanced to the end of the block that is read.
- **ssize_t write(int fd,void *buf,size_t count)**
Writes the given number of bytes to the given file, starting at the cursor. Bytes are written sequentially and the cursor is advanced to the end of the written block. Bytes to be written are located in the given buffer.
- **off_t lseek(int fd,off_t offset,int anchor)**
Move the cursor of the given file to the given location, based on the offset and anchor values.

Look up these functions; I recommend emulating this set of functions, as they will simplify the process of accessing the proper bytes from the filesystem.

You should also write two functions:

- **void displayBufferPage(uint8_t *buf,uint32_t count, uint32_t skip,uint64_t offset)**
Displays up to 256 bytes from the given buffer.
Parameters:
 - **buf** — buffer to be displayed
 - **count** — number of bytes to be displayed, max of 256
 - **skip** — number of bytes to skip, starting at **buf** (see examples)
 - **offset** — offset to display in the header.

- **void displayBuffer(uint8_t *buf, uint32_t count, uint64_t offset)**
Displays **count** bytes in the given buffer. The **offset** parameter is used to display the location of the buffer within the disk structure. This should just call **displayBufferPage()** repeatedly to display all of the bytes in the buffer.

The functions should display each 256-byte block in two forms: hexadecimal and character. A byte with offset *off* within the 256-byte block should be displayed in hexadecimal if $skip \leq off < skip + count$ and should also be displayed as a character if it is printable (**isprint()** returns true).

These display functions are for your use in building your program. They aren't part of the final product, but they may prove helpful in debugging steps along the way.

Format the output to make it easily readable. Here are some examples; your code doesn't have to replicate this format exactly, this is just to give you an idea of what the output looks like.

►Example 1

This is the VDI header from a sample VDI file. It's a 400-byte structure found at the start of a VDI file. A 400-byte buffer was read and displayed using the function call

```
displayBuffer(headerBuf,400,0);
```

```

1  Offset: 0x0
2      00 01 02 03 04 05 06 07 08 09 0a 0b 0c 0d 0e 0f      0...4...8...c...
3      +-----+-----+
4  00|3c 3c 3c 20 4f 72 61 63 6c 65 20 56 4d 20 56 69|00|<<< Oracle VM Vi|
5  10|72 74 75 61 6c 42 6f 78 20 44 69 73 6b 20 49 6d|10|rtualBox Disk Im|
6  20|61 67 65 20 3e 3e 3e 0a 00 00 00 00 00 00 00 00|20|age >>>
7  30|00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00|30|
8  40|7f 10 da be 01 00 01 00 90 01 00 00 02 00 00 00|40|
9  50|00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00|50|
10 60|00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00|60|
11 70|00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00|70|
12 80|00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00|80|
13 90|00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00|90|
14 a0|00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00|a0|
15 b0|00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00|b0|
16 c0|00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00|c0|
17 d0|00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00|d0|
18 e0|00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00|e0|
19 f0|00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00|f0|
20      +-----+-----+
21
22 Offset: 0x100
23      00 01 02 03 04 05 06 07 08 09 0a 0b 0c 0d 0e 0f      0...4...8...c...
24      +-----+-----+
25 00|00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00|00|
26 10|00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00|10|
27 20|00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00|20|
28 30|00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00|30|
29 40|00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00|40|
30 50|00 00 00 00 00 00 00 10 00 00 00 20 00 00 00 00|50|
31 60|00 00 00 00 00 00 00 00 00 00 02 00 00 00 00 00|60|
32 70|00 00 00 08 00 00 00 00 00 00 10 00 00 00 00 00|70|
33 80|80 00 00 00 80 00 00 00 07 07 d8 e6 24 6b c1 46|80|$k F
34 90|
35 a0|
36 b0|
37 c0|
38 d0|
39 e0|
40 f0|
41      +-----+-----+

```

►Example 2

This is the first 256 bytes of the VDI file's translation map. There is a field in the VDI header that determines the map's location within the VDI file; in this case the map is $1\,048\,576 = (10\,000)_{16}$ bytes from the start. It was called via

```
displayBufferPage(pageBuf, sizeof(pageBuf), 0, f->header->mapOffset);
```

```

1 Offset: 0x100000
2   00 01 02 03 04 05 06 07 08 09 0a 0b 0c 0d 0e 0f   0...4...8...c...
3   +-----+-----+-----+-----+-----+-----+
4 00|00 00 00 00 01 00 00 00 02 00 00 00 03 00 00 00|00|
5 10|04 00 00 00 05 00 00 00 06 00 00 00 07 00 00 00|10|
6 20|08 00 00 00 09 00 00 00 0a 00 00 00 0b 00 00 00|20|
7 30|0c 00 00 00 0d 00 00 00 0e 00 00 00 0f 00 00 00|30|
8 40|10 00 00 00 11 00 00 00 12 00 00 00 13 00 00 00|40|
9 50|14 00 00 00 15 00 00 00 16 00 00 00 17 00 00 00|50|
10 60|18 00 00 00 19 00 00 00 1a 00 00 00 1b 00 00 00|60|
11 70|1c 00 00 00 1d 00 00 00 1e 00 00 00 1f 00 00 00|70|
12 80|20 00 00 00 21 00 00 00 22 00 00 00 23 00 00 00|80|   !   "   #
13 90|24 00 00 00 25 00 00 00 26 00 00 00 27 00 00 00|90|$   %   &   '
14 a0|28 00 00 00 29 00 00 00 2a 00 00 00 2b 00 00 00|a0|(   )   *   +
15 b0|2c 00 00 00 2d 00 00 00 2e 00 00 00 2f 00 00 00|b0|,   -   .   /
16 c0|30 00 00 00 31 00 00 00 32 00 00 00 33 00 00 00|c0|0   1   2   3
17 d0|34 00 00 00 35 00 00 00 36 00 00 00 37 00 00 00|d0|4   5   6   7
18 e0|38 00 00 00 39 00 00 00 3a 00 00 00 3b 00 00 00|e0|8   9   :   ;
19 f0|3c 00 00 00 3d 00 00 00 3e 00 00 00 3f 00 00 00|f0|<  =  >  ?
20   +-----+-----+-----+-----+-----+-----+

```

►Example 3

The third example shows the partition table from the sample VDI file. The table — at least for MBR-based disks — is always in the first 512-byte sector at offset 446, or 190 bytes into the second 256-byte block. This example illustrates only showing bytes from *skip* through *skip + count*, as the first 190 bytes are skipped as are the last two. The function call used was

```
displayBufferPage(pageBuf,64,190,256);
```

```

1 Offset: 0x100
2   00 01 02 03 04 05 06 07 08 09 0a 0b 0c 0d 0e 0f   0...4...8...c...
3   +-----+-----+
4 00|                                     |00|
5 10|                                     |10|
6 20|                                     |20|
7 30|                                     |30|
8 40|                                     |40|
9 50|                                     |50|
10 60|                                     |60|
11 70|                                     |70|
12 80|                                     |80|
13 90|                                     |90|
14 a0|                                     |a0|
15 b0|                                     |b0|
16 c0|21 00 83 51 01 10 00 08 00 00 00 00 f8 03 00 00 00 00|c0|! Q
17 d0|00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00|d0|
18 e0|00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00|e0|
19 f0|00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00|f0|
20   +-----+-----+

```
