

Operating Systems

CSCI 5806

Spring Semester 2021 — CRN 21176

Term Project — Step 1 — VDI File Access

Target completion date: Sunday, January 31, 2021

Goals

- Provide the five basic file I/O functions to access disk space inside a VDI file.
- Create a structure or class to contain the data necessary to implement the five functions.

Details

You'll want a single entity — a structure or a class, either works — to represent a VDI file within your project. The intent is to collect the various data your project is going to need into one place for ease of use; if you're using C++ (you *are* using C or C++, right?) then a class can also contain the basic I/O functions as methods.

What do you need to keep track of in a VDI file? To implement the five functions you'll probably want at least these four items to be contained in your structure:

- The file descriptor for the VDI file. This is an `int`.
- A VDI header structure. Discussion of the structure can be found at <https://forums.virtualbox.org/viewtopic.php?t=8046>; C structures for the header can be found in the VirtualBox source code.
- The VDI translation map. This is optional; include it if you are going to enable dynamic VDI files (they're easy!), ignore it otherwise. This is just an array of integers, although you don't know the size in advance. The size is given in the VDI header, so you'll need to allocate this dynamically.
- A cursor. This is just an integer (of type `size_t`) that holds the location of the next byte to be read or written.

Once the structure / class is created, you'll need to implement a VDI version of each of the five basic I/O functions.

- **`struct VDIFile *vdiOpen(char *fn)`**
Open the file whose name is given. The filename can be either a relative or absolute path. The function returns a pointer to a dynamically created VDI file structure (see above), or a null pointer if there was an error. The function should load the header and the translation map, set the cursor to 0 and set the file descriptor to whatever was returned from the `open ()` system call.

If you're using a class, then this can return a boolean to indicate success or failure of the open.
- **`void vdiClose(struct VDIFile *f)`**
Close the file whose pointer is given. Deallocate any dynamically created memory regions.

- **ssize_t vdiRead(struct VDIFile *f,void *buf,size_t count)**
Reads the given number of bytes from the given VDI file's disk space, placing the bytes in the given buffer. The location of the first byte read is given by the cursor, which is relative to the start of the VDI file's data space. Advance the cursor by the number of bytes read.
- **ssize_t vdiWrite(struct VDIFile *f,void *buf,size_t count)**
Writes the given number of bytes to the given file, starting at the cursor. Bytes are written sequentially and the cursor is advanced to the end of the written block. Bytes to be written are located in the given buffer.
- **off_t vdiSeek(VDIFile *f,off_t offset,int anchor)**
Move the cursor of the given file to the given location, based on the offset and anchor values. If the resulting location is negative or larger than the disk size, do not change the value of the cursor.

If you are using a class, then the `VDIFile *` parameter is omitted.

The functions present the raw disk image as if it were the only data in the VDI file; all aspects of the VDI structure are to be hidden from the user. So, when the user wants to read, they only read from the disk space, never from the raw VDI file. Writing is similar. The cursor can only specify locations within the disk image.

If your project handles dynamic VDI files, then disk image should be presented as if all pages were laid out in the proper order, regardless of where they are in the VDI file. Missing pages must also be handled as if they were there.

You should also write a function that takes a pointer to a `VDIFile` as a parameter and displays its header fields in an easy-to-read manner. See example 1 below for a sample; your exact format may vary.

► Suggestions

- `vdiSeek()` should only set the cursor in the `VDIFile` structure; it should *not* call `lseek()`.
- Reading and writing should be done one page at a time. While there are bytes left to be read or written, do the following tasks:
 1. Determine where within the current page reading or writing should begin.
 2. Determine how many bytes should be read or written in the current page.
 3. Determine the physical location of the page. This may involve page translation and/or page allocation. Questions: What if the page is not allocated? What if the page is marked as "all zeroes"?
 4. Use `lseek()` to go to the proper location within the physical page.
 5. Use `read()` or `write()` to read or write *only* the bytes within the current page.
 6. Advance the cursor by the number of bytes read or written; subtract the number of bytes read or written from the number of bytes remaining.
- If you are planning to write into the filesystem, consider using the `mmap()` and `munmap()` functions to load the VDI file header and translation map. These act like a write-through cache; they read the areas from the file into memory set up by the OS, and writing to them automatically writes back to the file.

►Example 1

This is the output from my **testVDI()** function on the good-fixed-1k.vdi file.

```

1 Dump of VDI header:
2   Image name: [<<< Oracle VM VirtualBox Disk Image >>>
3               ]
4   Signature: 0xbeda107f
5   Version: 1.01
6   Header size: 0x000000190 400
7   Image type: 0x000000002
8   Flags: 0x000000000
9   Virtual CHS: 0-0-0
10  Sector size: 0x000000200 512
11  Logical CHS: 260-16-63
12  Sector size: 0x000000200 512
13  Map offset: 0x00100000 1048576
14  Frame offset: 0x00200000 2097152
15  Frame size: 0x00100000 1048576
16 Extra frame size: 0x000000000 0
17  Total frames: 0x000000080 128
18 Frames allocated: 0x000000080 128
19  Disk size: 0x00000000008000000 134217728
20  UUID: 6fd7a9b3-5226-8740-aea7-506076e113b0
21  Last snap UUID: 6221a1ea-f266-e946-aeeb-4959370a749b
22  Link UUID: 00000000-0000-0000-0000-000000000000
23  Parent UUID: 00000000-0000-0000-0000-000000000000
24  Image comment:
25 Offset: 0x54
26   00 01 02 03 04 05 06 07 08 09 0a 0b 0c 0d 0e 0f  0...4...8...c...
27   +-----+-----+
28 00|00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00|00|
29 10|00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00|10|
30 20|00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00|20|
31 30|00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00|30|
32 40|00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00|40|
33 50|00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00|50|
34 60|00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00|60|
35 70|00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00|70|
36 80|00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00|80|
37 90|00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00|90|
38 a0|00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00|a0|
39 b0|00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00|b0|
40 c0|00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00|c0|
41 d0|00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00|d0|
42 e0|00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00|e0|
43 f0|00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00|f0|
44   +-----+-----+
45
46 Unused: 0

```

49 Header in binary:

50 Offset: 0x0

```

51      00 01 02 03 04 05 06 07 08 09 0a 0b 0c 0d 0e 0f      0...4...8...c...
52      +-----+-----+-----+-----+-----+-----+
53 00|3c 3c 3c 20 4f 72 61 63 6c 65 20 56 4d 20 56 69|00|<<< Oracle VM Vi|
54 10|72 74 75 61 6c 42 6f 78 20 44 69 73 6b 20 49 6d|10|rtualBox Disk Im|
55 20|61 67 65 20 3e 3e 3e 0a 00 00 00 00 00 00 00 00|20|age >>>|
56 30|00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00|30|
57 40|7f 10 da be 01 00 01 00 90 01 00 00 02 00 00 00 00|40|
58 50|00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00|50|
59 60|00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00|60|
60 70|00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00|70|
61 80|00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00|80|
62 90|00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00|90|
63 a0|00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00|a0|
64 b0|00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00|b0|
65 c0|00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00|c0|
66 d0|00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00|d0|
67 e0|00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00|e0|
68 f0|00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00|f0|
69      +-----+-----+-----+-----+-----+-----+

```

70

71 Offset: 0x100

```

72      00 01 02 03 04 05 06 07 08 09 0a 0b 0c 0d 0e 0f      0...4...8...c...
73      +-----+-----+-----+-----+-----+-----+
74 00|00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00|00|
75 10|00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00|10|
76 20|00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00|20|
77 30|00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00|30|
78 40|00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00|40|
79 50|00 00 00 00 00 00 00 10 00 00 00 20 00 00 00 00 00|50|
80 60|00 00 00 00 00 00 00 00 00 00 02 00 00 00 00 00 00|60|
81 70|00 00 00 08 00 00 00 00 00 00 00 10 00 00 00 00 00|70|
82 80|80 00 00 00 80 00 00 00 6f d7 a9 b3 52 26 87 40|80|          o   R& @
83 90|                                     |90|
84 a0|                                     |a0|
85 b0|                                     |b0|
86 c0|                                     |c0|
87 d0|                                     |d0|
88 e0|                                     |e0|
89 f0|                                     |f0|
90      +-----+-----+-----+-----+-----+-----+

```

92	Translation map:																								
93	Offset: 0x0																								
94	00 01 02 03 04 05 06 07 08 09 0a 0b 0c 0d 0e 0f																0...4...8...c...								
95	+-----+ +-----+																+-----+								
96	00	00	00	00	00	01	00	00	00	02	00	00	00	03	00	00	00	00							
97	10	04	00	00	00	05	00	00	00	06	00	00	00	07	00	00	00	10							
98	20	08	00	00	00	09	00	00	00	0a	00	00	00	0b	00	00	00	20							
99	30	0c	00	00	00	0d	00	00	00	0e	00	00	00	0f	00	00	00	30							
100	40	10	00	00	00	11	00	00	00	12	00	00	00	13	00	00	00	40							
101	50	14	00	00	00	15	00	00	00	16	00	00	00	17	00	00	00	50							
102	60	18	00	00	00	19	00	00	00	1a	00	00	00	1b	00	00	00	60							
103	70	1c	00	00	00	1d	00	00	00	1e	00	00	00	1f	00	00	00	70							
104	80	20	00	00	00	21	00	00	00	22	00	00	00	23	00	00	00	80	!	"	#				
105	90	24	00	00	00	25	00	00	00	26	00	00	00	27	00	00	00	90	\$	%	&	'			
106	a0	28	00	00	00	29	00	00	00	2a	00	00	00	2b	00	00	00	a0	()	*	+			
107	b0	2c	00	00	00	2d	00	00	00	2e	00	00	00	2f	00	00	00	b0	,	-	.	/			
108	c0	30	00	00	00	31	00	00	00	32	00	00	00	33	00	00	00	c0	0	1	2	3			
109	d0	34	00	00	00	35	00	00	00	36	00	00	00	37	00	00	00	d0	4	5	6	7			
110	e0	38	00	00	00	39	00	00	00	3a	00	00	00	3b	00	00	00	e0	8	9	:	;			
111	f0	3c	00	00	00	3d	00	00	00	3e	00	00	00	3f	00	00	00	f0	<	=	>	?			
112	+-----+ +-----+																+-----+								
113																									
114	Offset: 0x100																								
115	00 01 02 03 04 05 06 07 08 09 0a 0b 0c 0d 0e 0f																0...4...8...c...								
116	+-----+ +-----+																+-----+								
117	00	40	00	00	00	41	00	00	00	42	00	00	00	43	00	00	00	00	@	A	B	C			
118	10	44	00	00	00	45	00	00	00	46	00	00	00	47	00	00	00	10	D	E	F	G			
119	20	48	00	00	00	49	00	00	00	4a	00	00	00	4b	00	00	00	20	H	I	J	K			
120	30	4c	00	00	00	4d	00	00	00	4e	00	00	00	4f	00	00	00	30	L	M	N	O			
121	40	50	00	00	00	51	00	00	00	52	00	00	00	53	00	00	00	40	P	Q	R	S			
122	50	54	00	00	00	55	00	00	00	56	00	00	00	57	00	00	00	50	T	U	V	W			
123	60	58	00	00	00	59	00	00	00	5a	00	00	00	5b	00	00	00	60	X	Y	Z	[
124	70	5c	00	00	00	5d	00	00	00	5e	00	00	00	5f	00	00	00	70	\]	^	_			
125	80	60	00	00	00	61	00	00	00	62	00	00	00	63	00	00	00	80	'	a	b	c			
126	90	64	00	00	00	65	00	00	00	66	00	00	00	67	00	00	00	90	d	e	f	g			
127	a0	68	00	00	00	69	00	00	00	6a	00	00	00	6b	00	00	00	a0	h	i	j	k			
128	b0	6c	00	00	00	6d	00	00	00	6e	00	00	00	6f	00	00	00	b0	l	m	n	o			
129	c0	70	00	00	00	71	00	00	00	72	00	00	00	73	00	00	00	c0	p	q	r	s			
130	d0	74	00	00	00	75	00	00	00	76	00	00	00	77	00	00	00	d0	t	u	v	w			
131	e0	78	00	00	00	79	00	00	00	7a	00	00	00	7b	00	00	00	e0	x	y	z	{			
132	f0	7c	00	00	00	7d	00	00	00	7e	00	00	00	7f	00	00	00	f0		}	~				
133	+-----+ +-----+																+-----+								

```

137 Partition table from Master Boot Record:
138 Offset: 0x100
139      00 01 02 03 04 05 06 07 08 09 0a 0b 0c 0d 0e 0f      0...4...8...c...
140      +-----+-----+
141 00|                                     |00|
142 10|                                     |10|
143 20|                                     |20|
144 30|                                     |30|
145 40|                                     |40|
146 50|                                     |50|
147 60|                                     |60|
148 70|                                     |70|
149 80|                                     |80|
150 90|                                     |90|
151 a0|                                     |a0|
152 b0|                                     |b0|
153 c0|01 04 83 01 82 02 00 08 00 00 00 f8 03 00 00 00|c0|
154 d0|00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00|d0|
155 e0|00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00|e0|
156 f0|00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00|f0|
157      +-----+-----+

```

►Other Examples

The output for the other five sample VDI files is available in the repository in the file `step1.log`.